

R PROGRAMMING

An Introduction





INTRODUCTIONS





HELLO
my name is

Brad

-  bradleyboehmke.github.io
-  bradleyboehmke@gmail.com
-  [@bradleyboehmke](https://twitter.com/bradleyboehmke)
-  [bradleyboehmke](https://github.com/bradleyboehmke)

COURSE OBJECTIVES

Provides an intensive, hands-on introduction to the R programming language.

- Get up and running with R & RStudio
- Understand the basics of data prep
- Perform data transformation
- Visualize data
- Compute basic statistical inference


SCHEDULE

8:00 - 9:00:	Intro
9:00 - 10:00:	Data Prep
10:00 - 10:20:	Break
10:20 - 12:00:	Data Transformation
12:00 - 1:30:	Lunch
1:30 - 2:45:	Visualization
2:45 - 3:00:	Break
3:00 - 4:30:	Interactive Tutorial

Lots of hands-on coding exercises



You will be overwhelmed!




Strong proponent of
collaborative work!



CLASS MATERIAL

<https://uc-r.github.io>


- Class material
- Tutorials
- Additional resources links
- Exercises



UNIVERSITY OF
Cincinnati

Follow us on twitter @UC_Rstats

- Home
- Course: Intro to R Bootcamp
- Course: Data Wrangling with R
- Introduction to R
- R Basics
- Workflow
 - RStudio Projects
 - R Markdown
 - R Notebook
- Data Types
 - Dealing with Numbers
 - Dealing with Characters
 - Dealing with Regex
 - Dealing with Factors

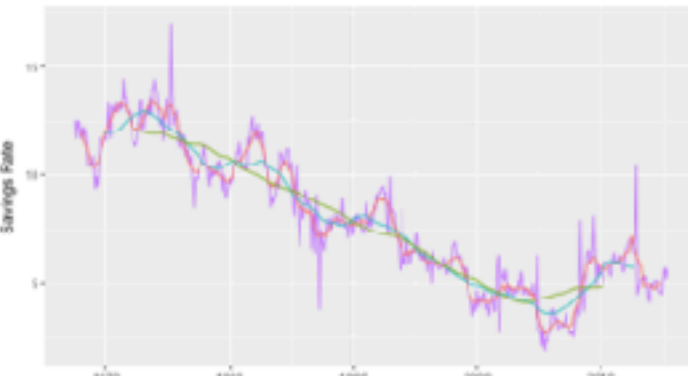


UC Business Analytics R Programming Guide

Moving Averages

04 Aug 2017

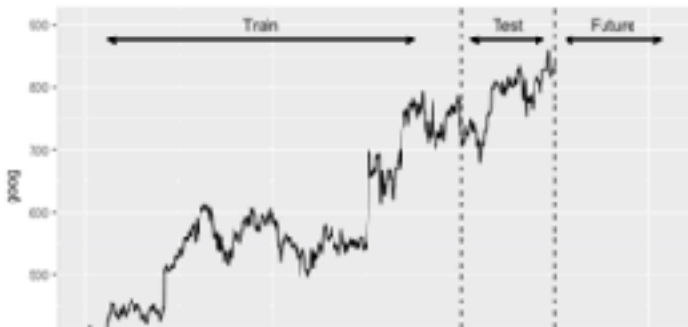
Smoothing methods are a family of forecasting approaches that average values over multiple periods in order to reduce the noise and uncover patterns in the data. Moving averages are one such smoothing method. In this new [time series tutorial](#), you will learn the basics of performing smoothing averages.



Benchmark Methods & Forecast Accuracy

16 Jun 2017

In this new [time series tutorial](#), you will learn general tools that are useful for many different forecasting situations. It will describe some methods for benchmark forecasting, methods for checking whether a forecasting model has



PROGRAMMING & ANALYSIS



Download and install R, a free software environment for statistical computing and graphics from CRAN, the Comprehensive R Archive Network. It is highly recommended to install a precompiled binary distribution for your operating system; follow these instructions:

1. Go to <https://cran.r-project.org/>
2. Click “Download R for Mac/Windows”
3. Download the appropriate file:
 - Windows users click Base, and download the installer for the latest R version
 - Mac users select the file R-3.X.X.pkg that aligns with your OS version
4. Follow the instructions of the installer.

PROGRAMMING & ANALYSIS



Install RStudio's IDE (stands for integrated development environment), a powerful user interface for R. RStudio includes a text editor, so you do not have to install another stand-alone editor. Follow these instructions:

1. Go to RStudio for desktop <https://www.rstudio.com/products/rstudio/download/>
2. Select the install file for your OS
3. Follow the instructions of the installer.

There are other R IDE's available: Emacs, Microsoft R Open, Notepad++, etc; however, I have found RStudio to be my preferred route. When you are done installing RStudio click on the icon.

QUESTIONS ABOUT THE CLASS?



FUNDAMENTALS



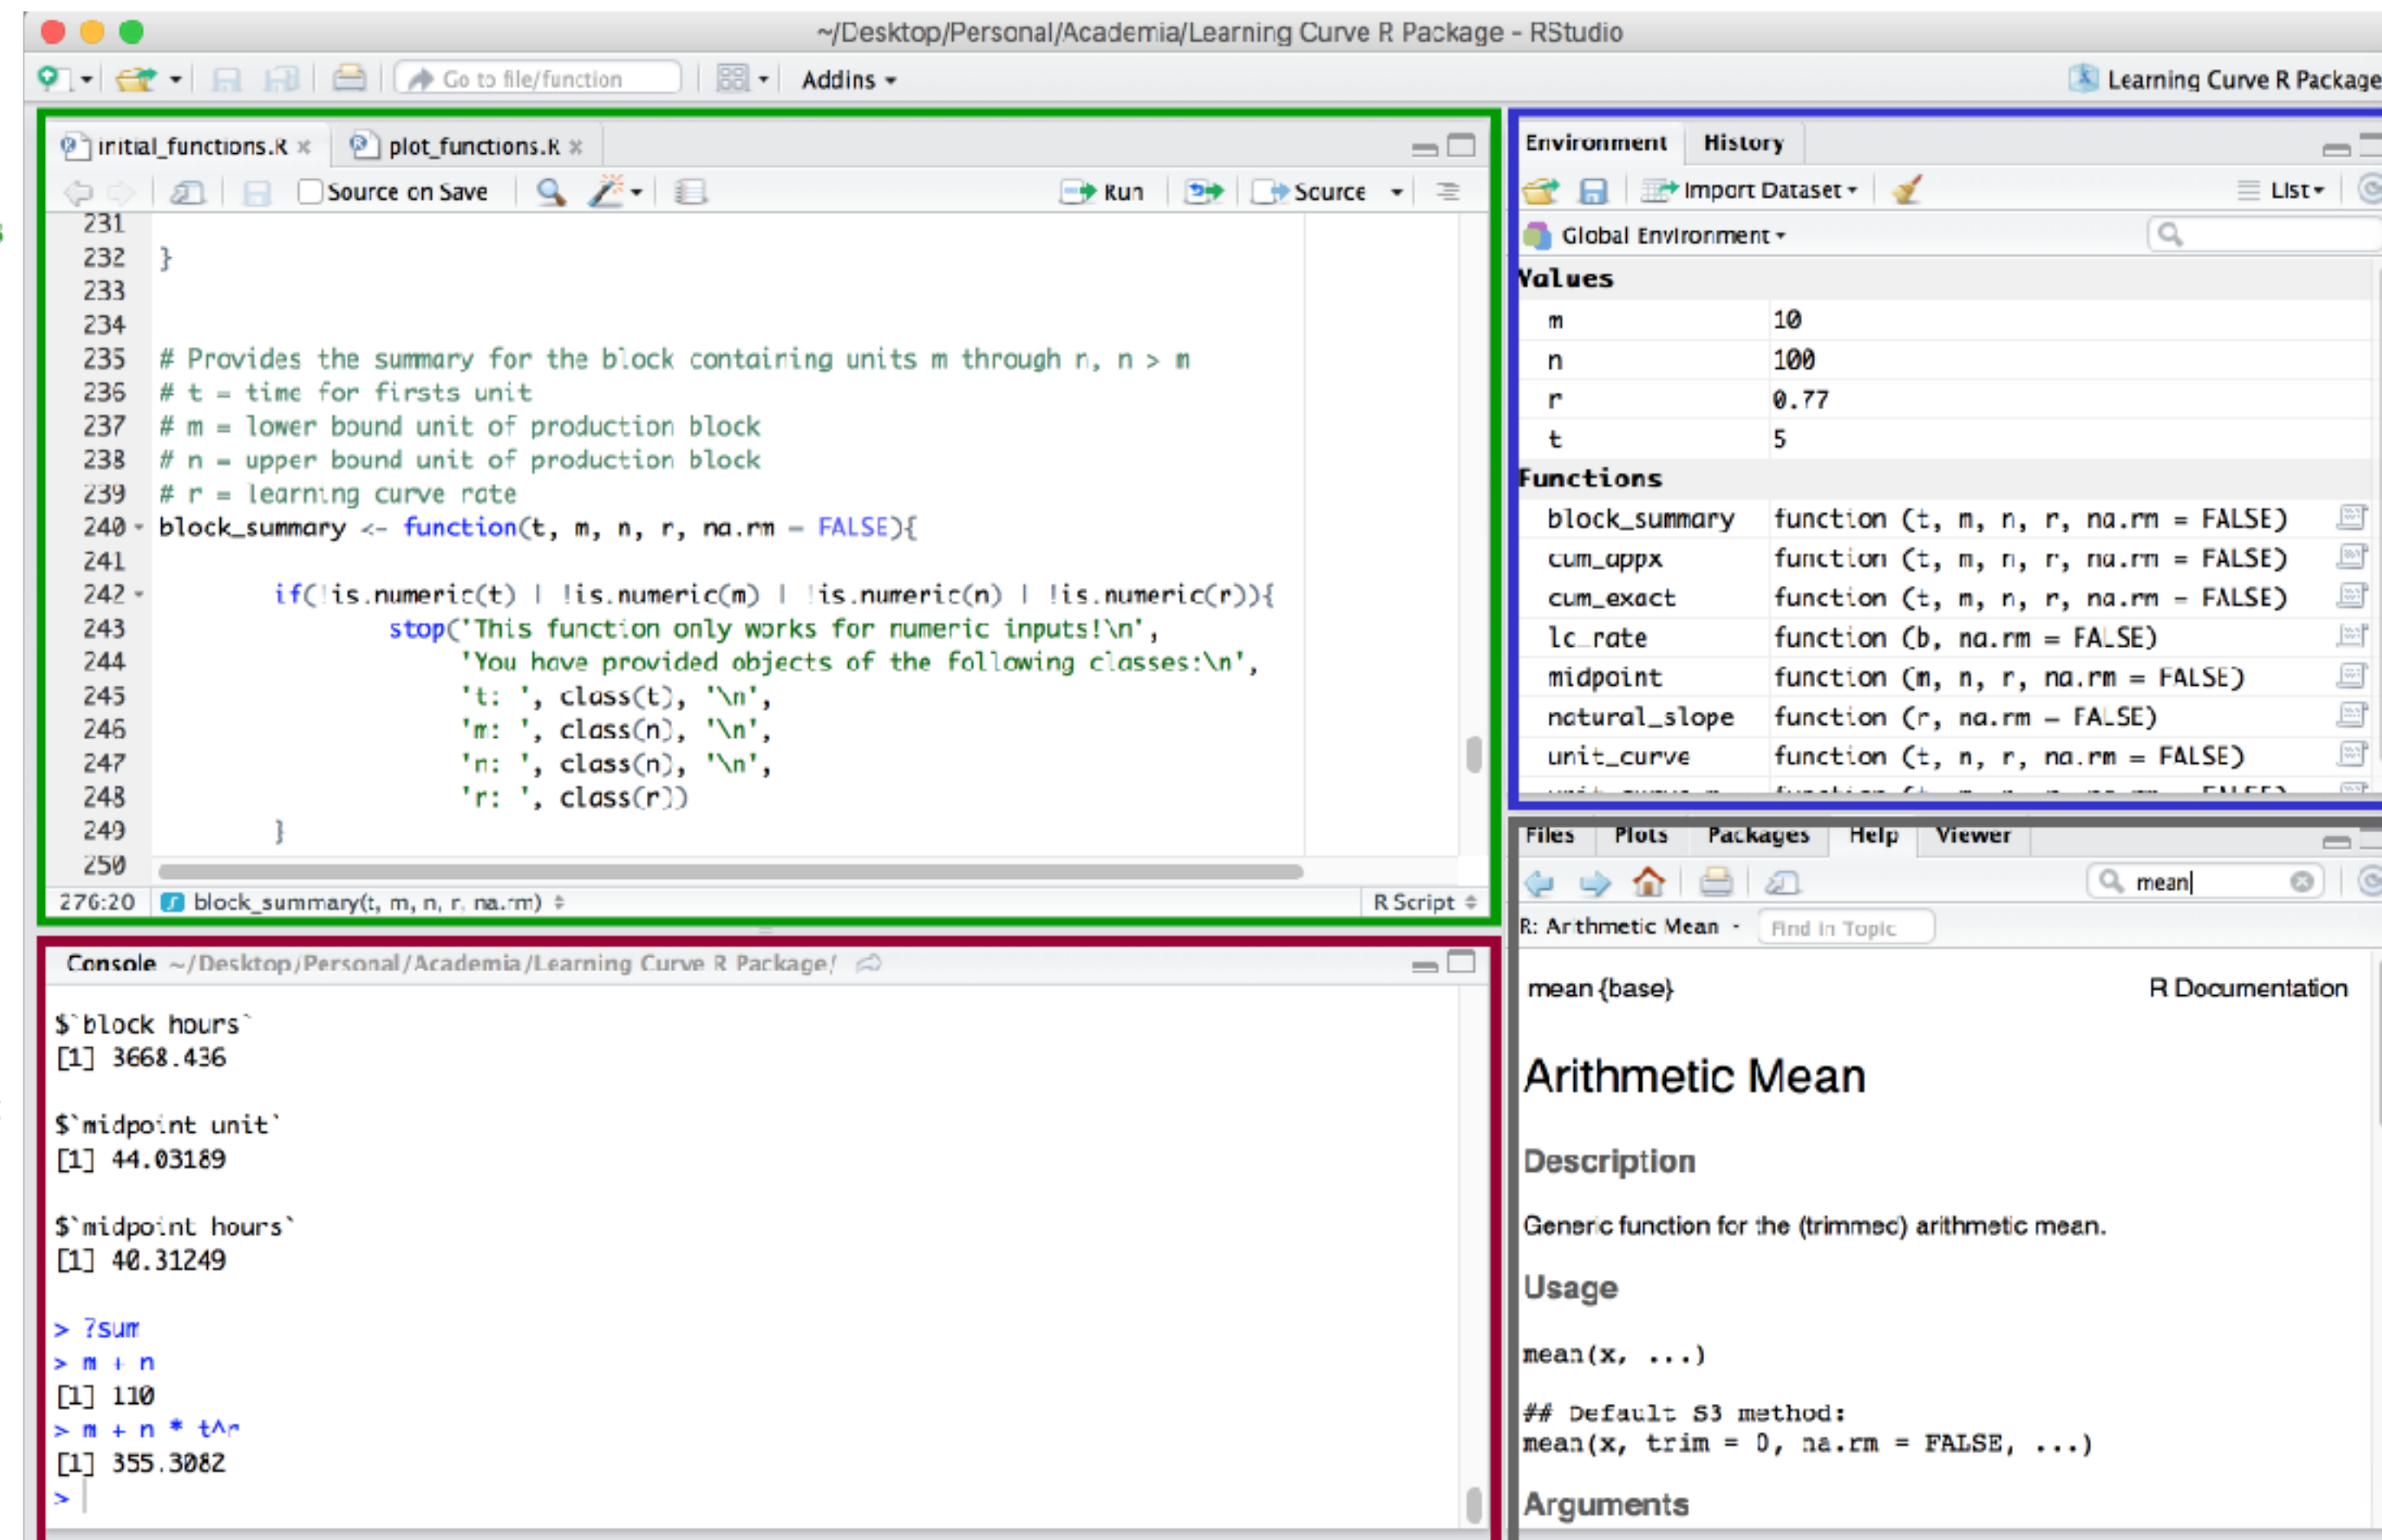
UNDERSTANDING THE RSTUDIO IDE

Script files

- Saves your script
- Allows code & comments
- Can have multiple files open at a time

Console/Command line

- Can use as calculator
- Does not save code
- This is where your output is displayed



Workspace environment

- Holds your objects
- Can review history

Misc - Displays:

- files in working directory
- plots when produced
- help files/search

Thorough tutorial regarding the RStudio console: <http://dss.princeton.edu/training/RStudio101.pdf>

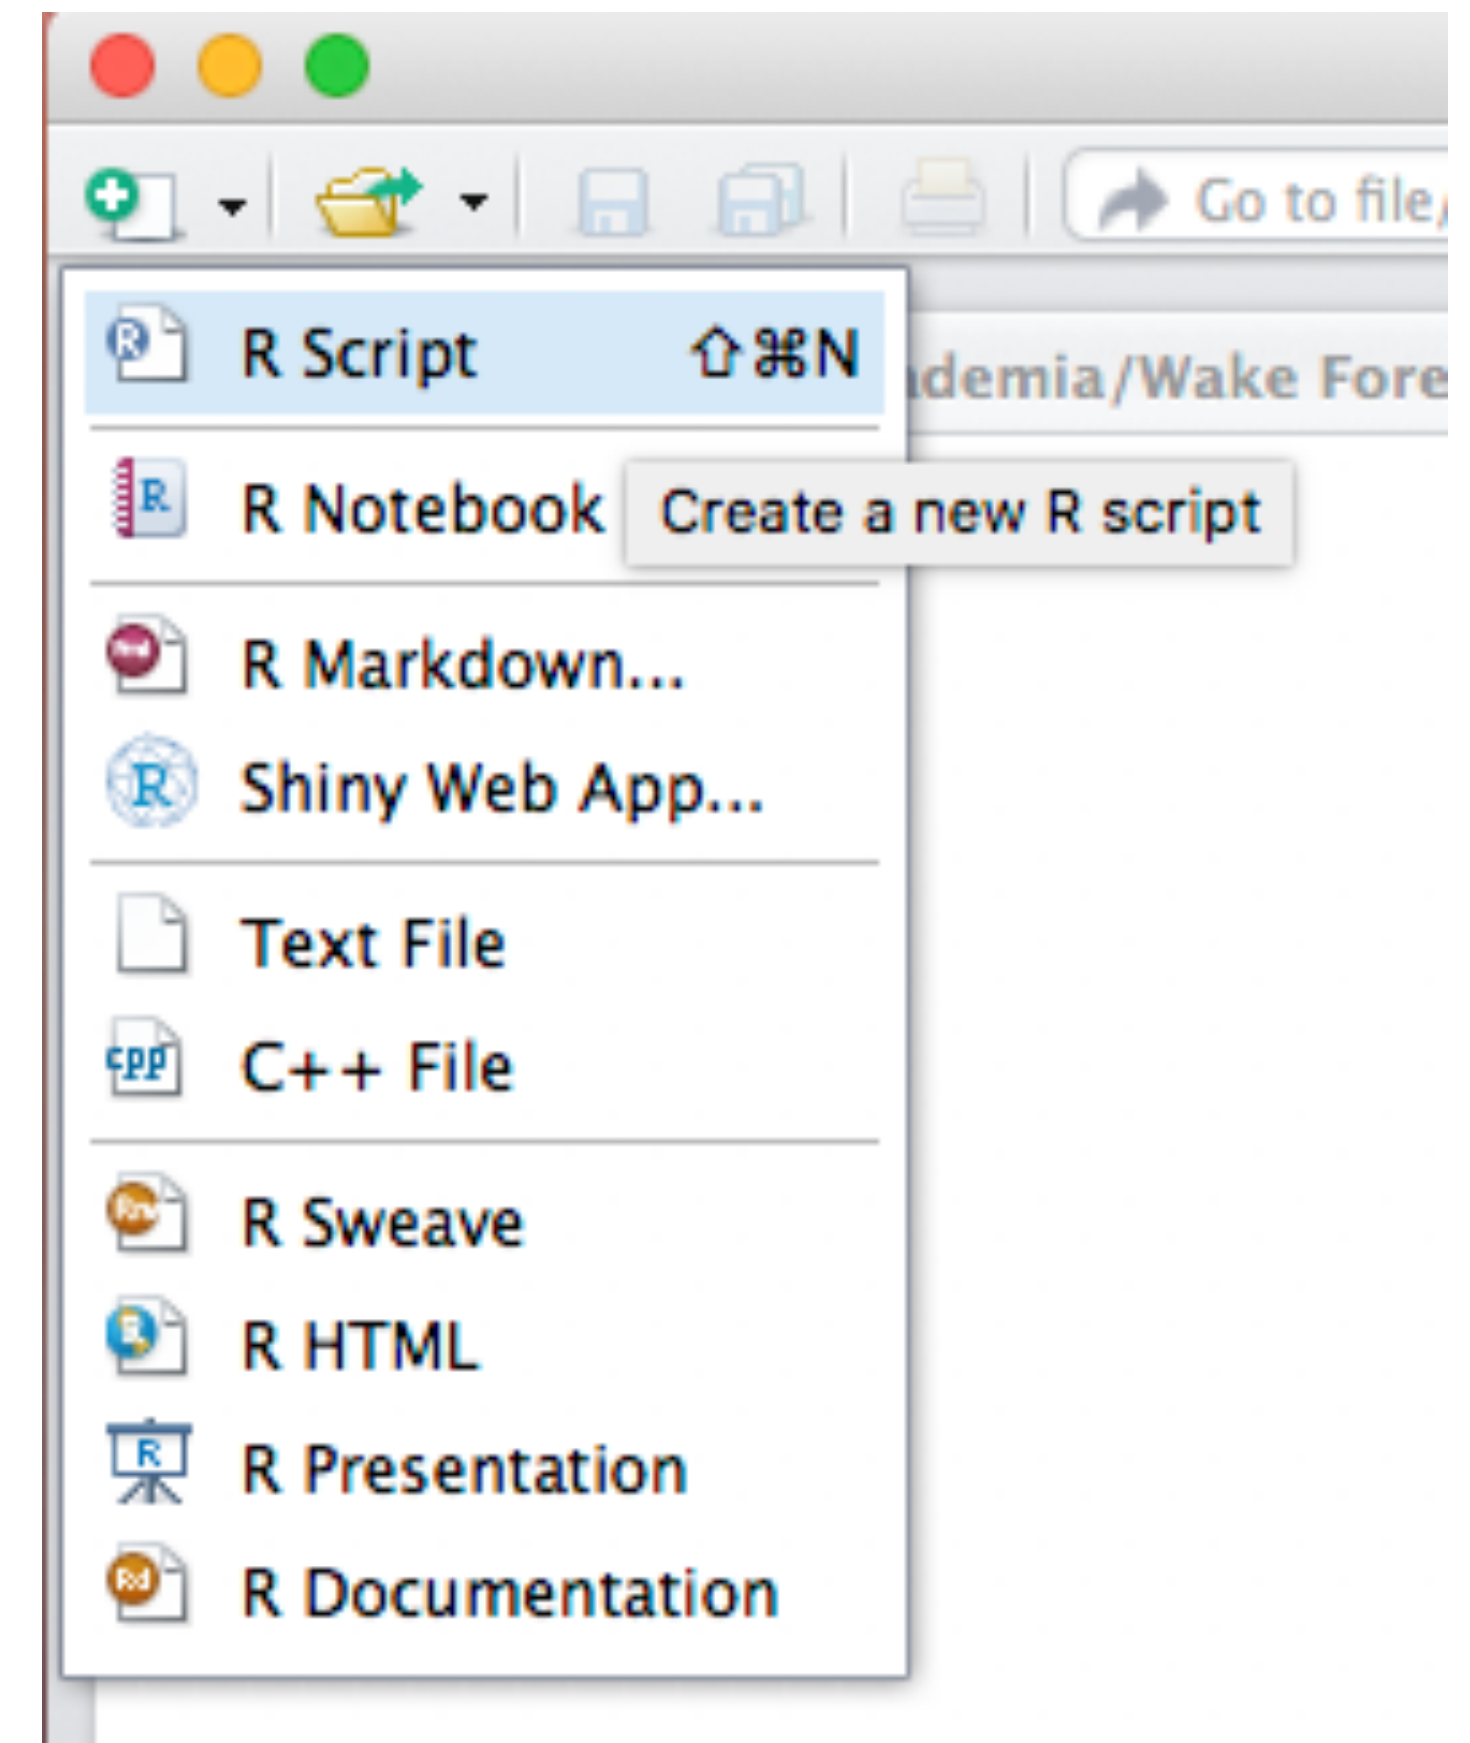
YOUR TURN!

1. *Create a new .R script named my_first_script.R*
2. *Write and execute the following code in the .R script and identify where in the IDE the outputs can be found.*

```
mtcars  
?sum  
hist(mtcars$mpg)  
avg_mpg <- mean(mtcars$mpg)  
random_numbers <- runif(25)  
history()
```

SOLUTION

1. *Create a new .R script*



SOLUTION

mtcars

?sum

hist(mtcars\$mpg)

avg_mpg <- mean(mtcars\$mpg)

random_numbers <- runif(25)

history()

The screenshot shows the RStudio interface with the following components:

- Script Editor:** Contains the following R code:

```
1 mtcars
2 ?sum
3 hist(mtcars$mpg)
4 avg_mpg <- mean(mtcars$mpg)
5 random_numbers <- runif(25)
6 history()
7
```
- Environment:** Shows "Global Environment" and "Environment is empty".
- Files:** Shows the project structure: Home > Dropbox > Academia > Wake Forest > BAN6003. The file "BAN6003.Rproj" is listed with a size of 205 B and a modification date of Jul 5, 2017, 11:15 AM.
- Console:** Shows the output of the R code:

```
> mtcars
      mpg  cyl  disp  hp  drat    wt  qsec vs  am  gear  carb
Mazda RX4         21.0   6  160.0 110  3.90 2.620 16.46  0  1    4    4
Mazda RX4 Wag     21.0   6  160.0 110  3.90 2.875 17.02  0  1    4    4
Datsun 710         22.8   4  108.0  93  3.85 2.320 18.61  1  1    4    1
Hornet 4 Drive     21.4   6  258.0 110  3.08 3.215 19.44  1  0    3    1
Hornet Sportabout 18.7   8  360.0 175  3.15 3.440 17.02  0  0    3    2
Valiant           18.1   6  225.0 105  2.76 3.460 20.22  1  0    3    1
Duster 360        14.3   8  360.0 245  3.21 3.570 15.84  0  0    3    4
Merc 240D          24.4   4  146.7  62  3.69 3.190 20.00  1  0    4    2
Merc 230           22.8   4  140.8  95  3.92 3.150 22.90  1  0    4    2
Merc 280           19.2   6  167.6 123  3.92 3.440 18.30  1  0    4    4
Merc 280C          17.8   6  167.6 123  3.92 3.440 18.90  1  0    4    4
Merc 450SE         16.4   8  275.8 180  3.07 4.070 17.40  0  0    3    3
Merc 450SL         17.3   8  275.8 180  3.07 3.730 17.60  0  0    3    3
Merc 450SLC        15.2   8  275.8 180  3.07 3.730 18.00  0  0    3    3
Cadillac Fleetwood 10.4   8  472.0 205  2.93 5.250 17.98  0  0    3    4
Lincoln Continental 10.4   8  460.0 215  3.00 5.424 17.82  0  0    3    4
Chrysler Imperial 14.7   8  440.0 230  3.23 5.345 17.42  0  0    3    4
Fiat 128           32.4   4   78.7  66  4.08 2.200 19.47  1  1    4    1
```

SOLUTION

mtcars

?sum

hist(mtcars\$mpg)

avg_mpg <- mean(mtcars\$mpg)

random_numbers <- runif(25)

history()

The screenshot shows the RStudio interface with the following components:

- Script Editor:** Contains a script named "my first script.R" with the following code:

```
1 mtcars
2 ?sum
3 hist(mtcars$mpg)
4 avg_mpg <- mean(mtcars$mpg)
5 random_numbers <- runif(25)
6 history()
7 |
```
- Console:** Displays the output of the script execution, showing a table of car data and the results of the `?sum` command.

Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0
Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0
Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0
Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1
Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1
Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1
Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0
Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0
AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0
Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0
Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0
Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1
Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1
Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1
Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1
Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1
Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1
Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1
- Environment:** Shows the "Global Environment" with the message "Environment is empty".
- Help Window:** Displays the documentation for the `sum` function, titled "Sum of Vector Elements". It includes a description, usage, arguments, and details.

SOLUTION

mtcars

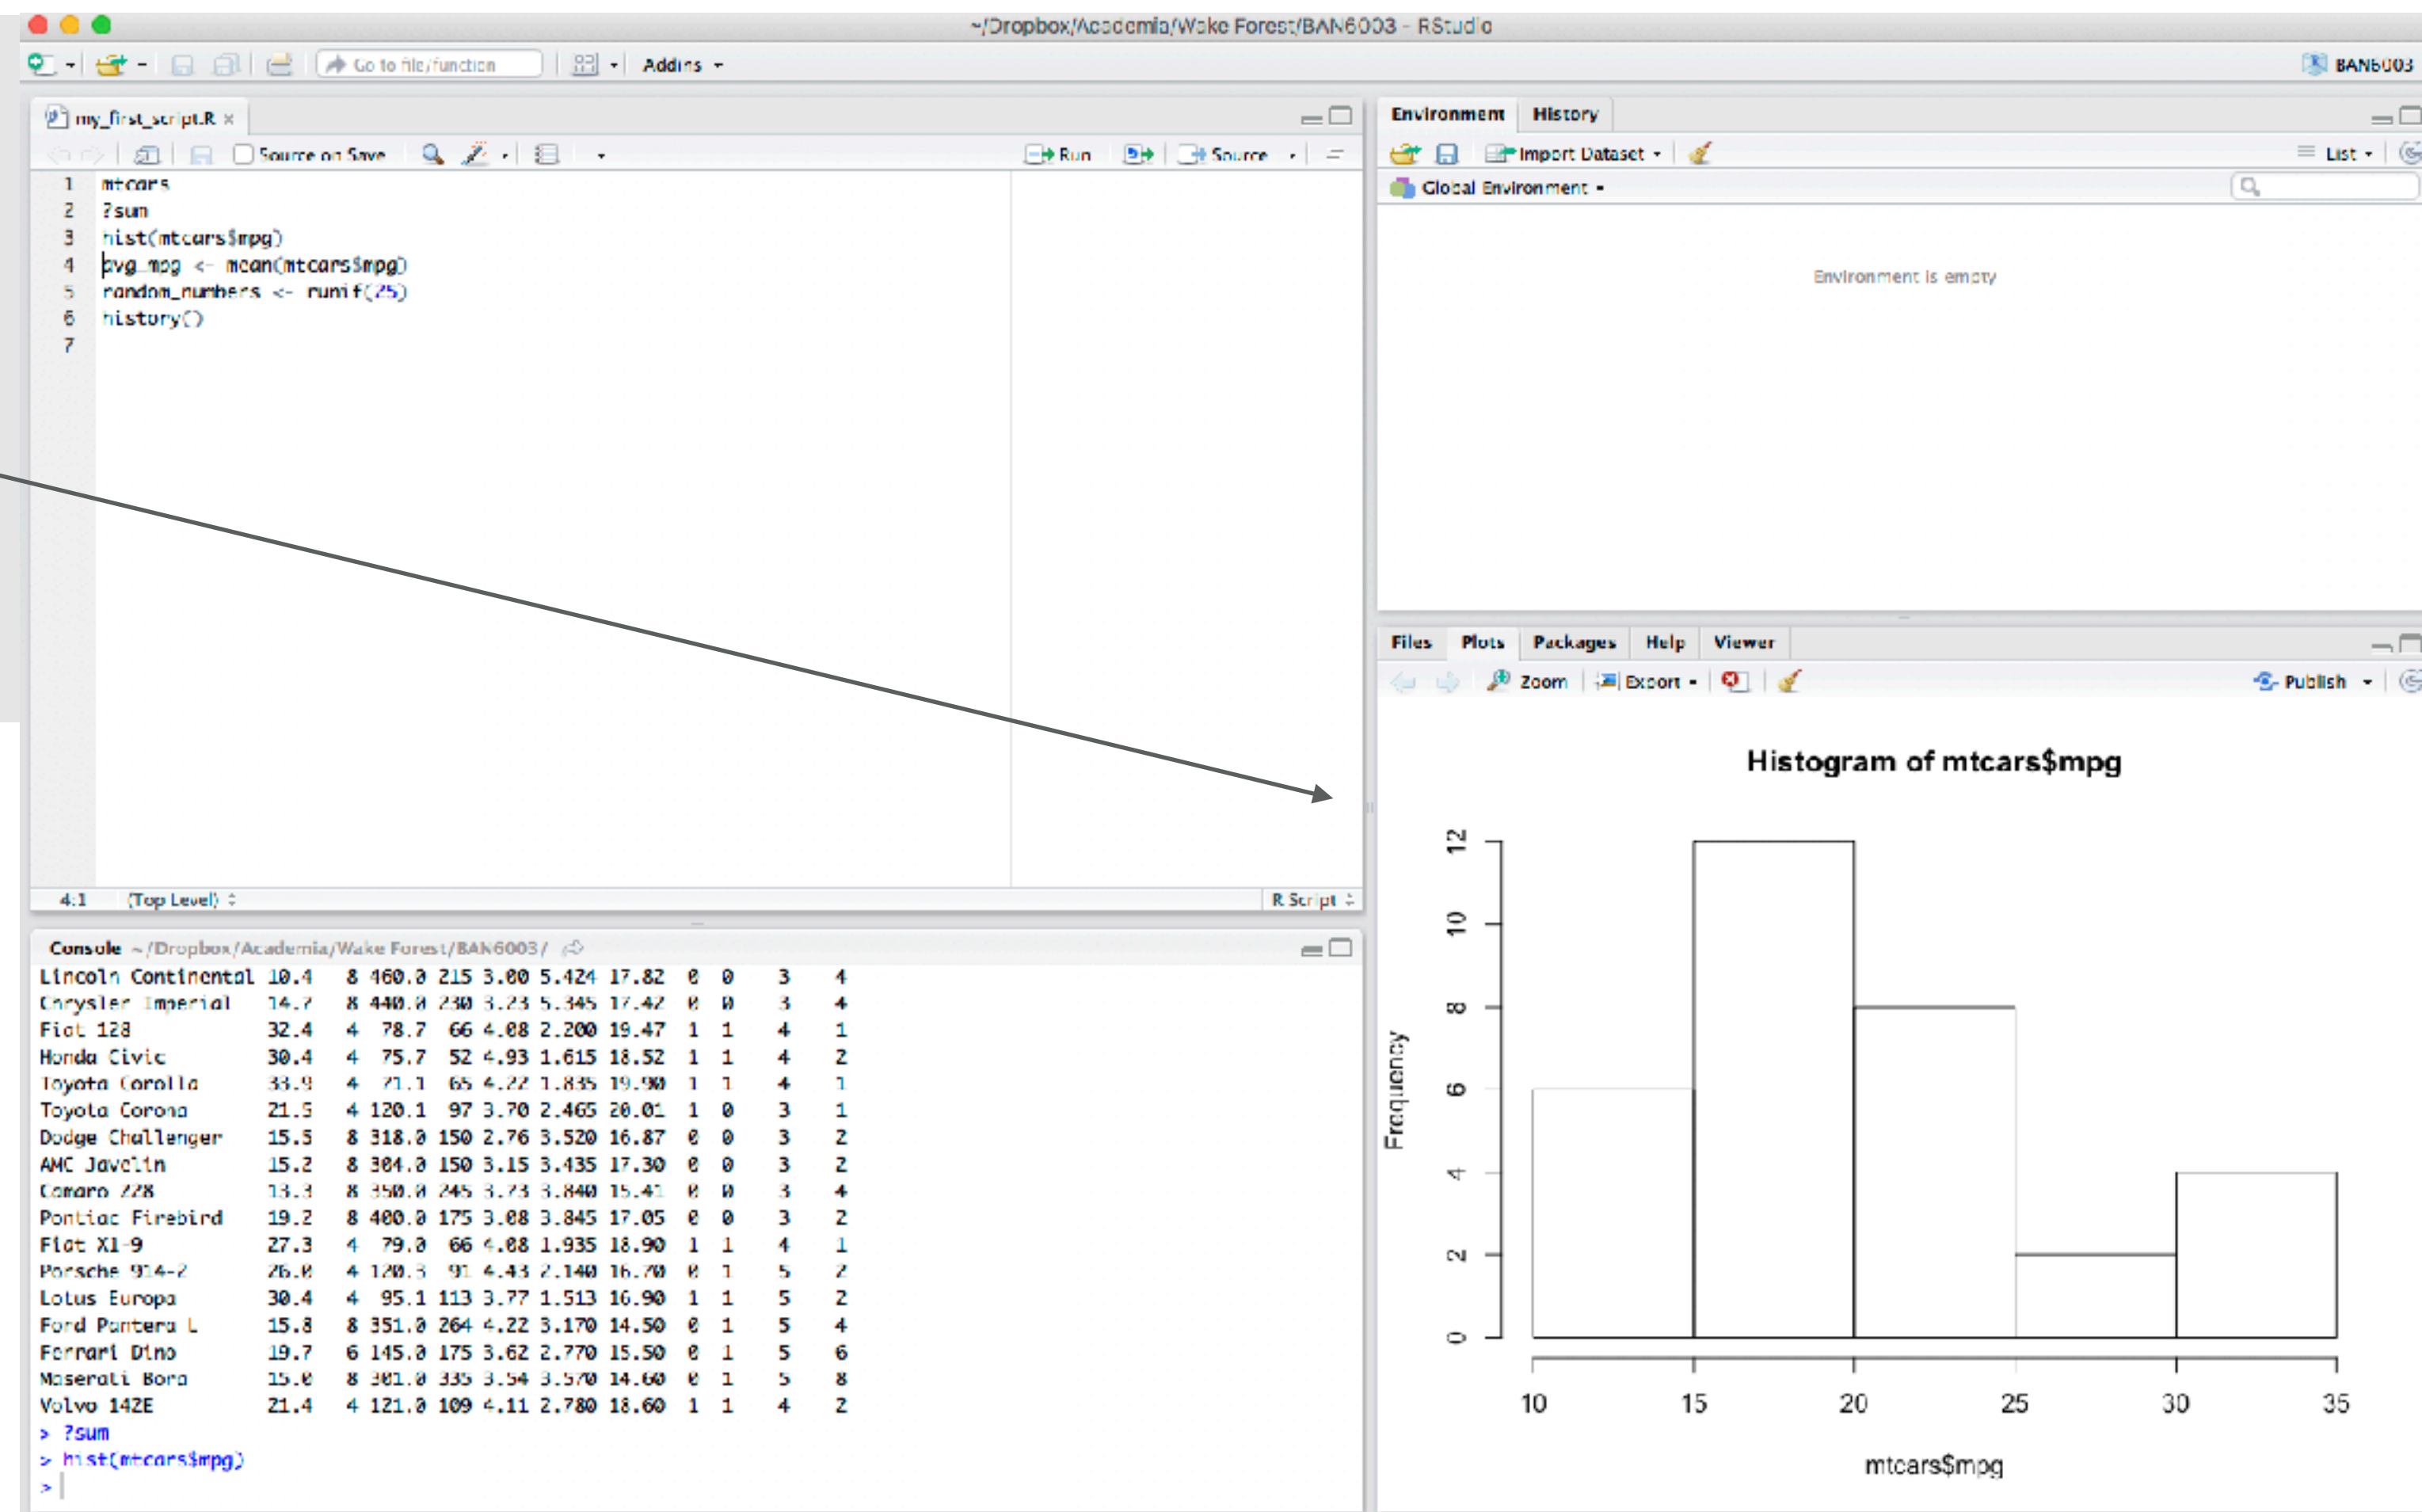
?sum

hist(mtcars\$mpg)

avg_mpg <- mean(mtcars\$mpg)

random_numbers <- runif(25)

history()



SOLUTION

mtcars

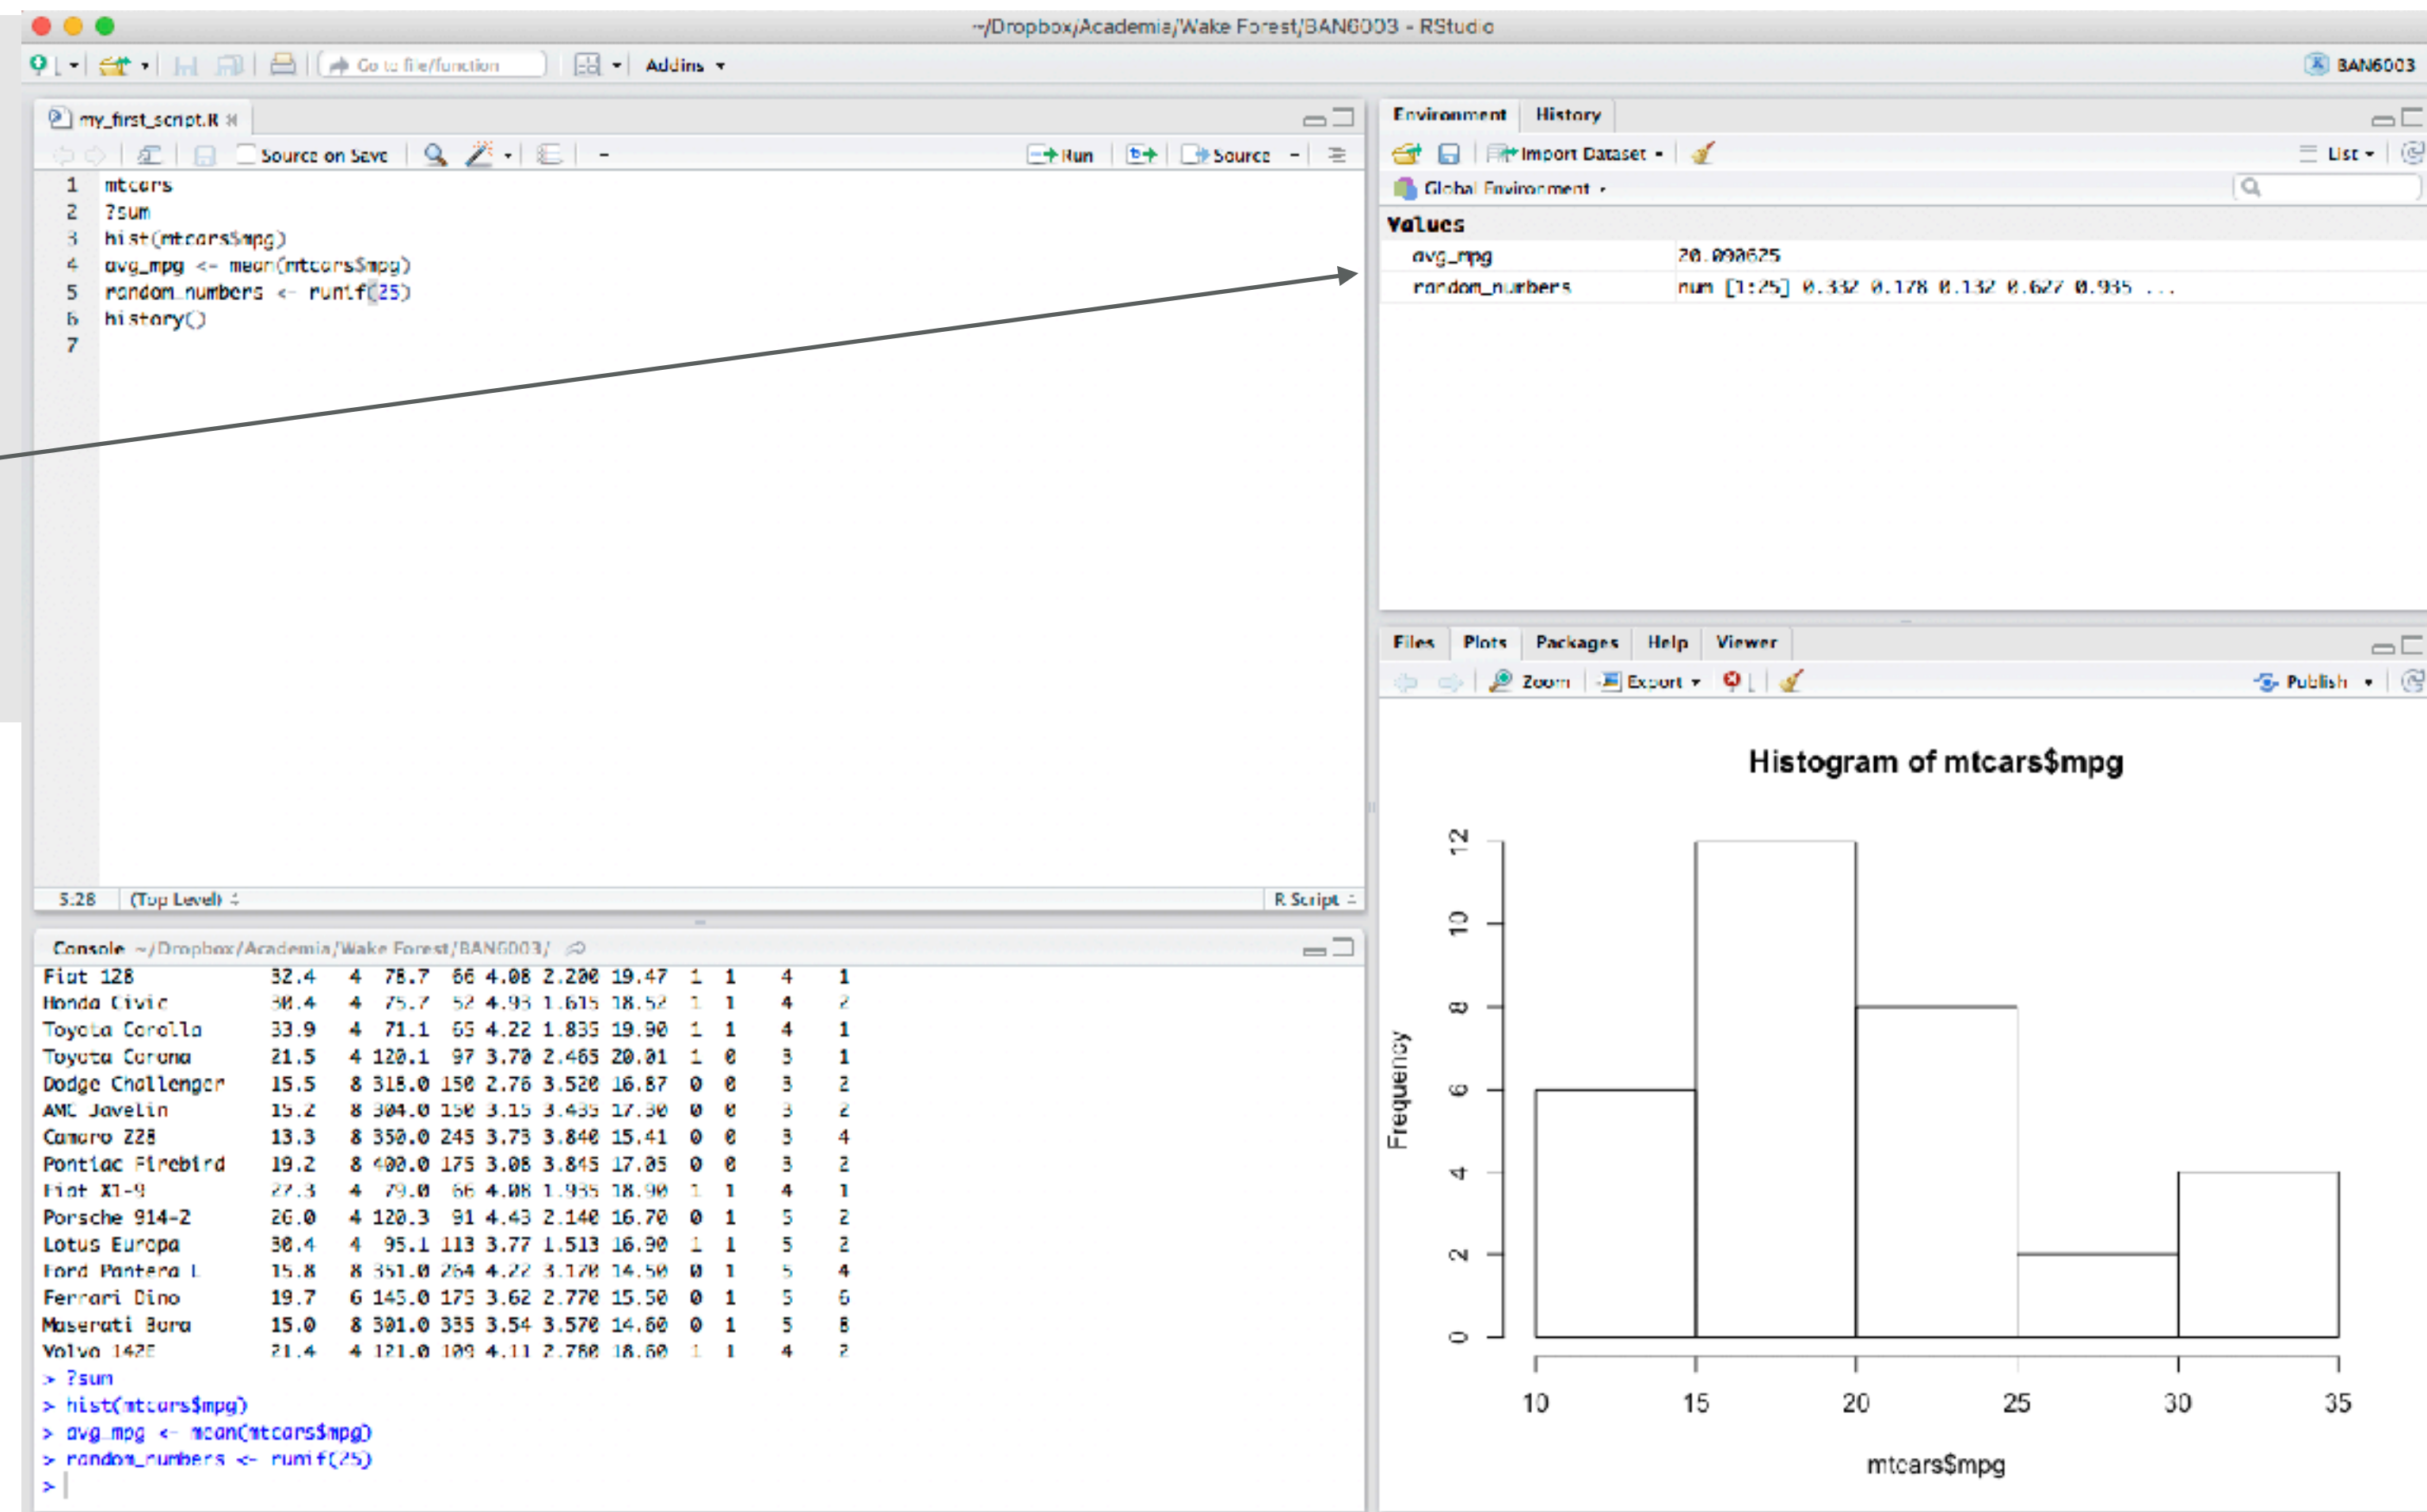
?sum

hist(mtcars\$mpg)

avg_mpg <- mean(mtcars\$mpg)

random_numbers <- runif(25)

history()



SOLUTION

mtcars

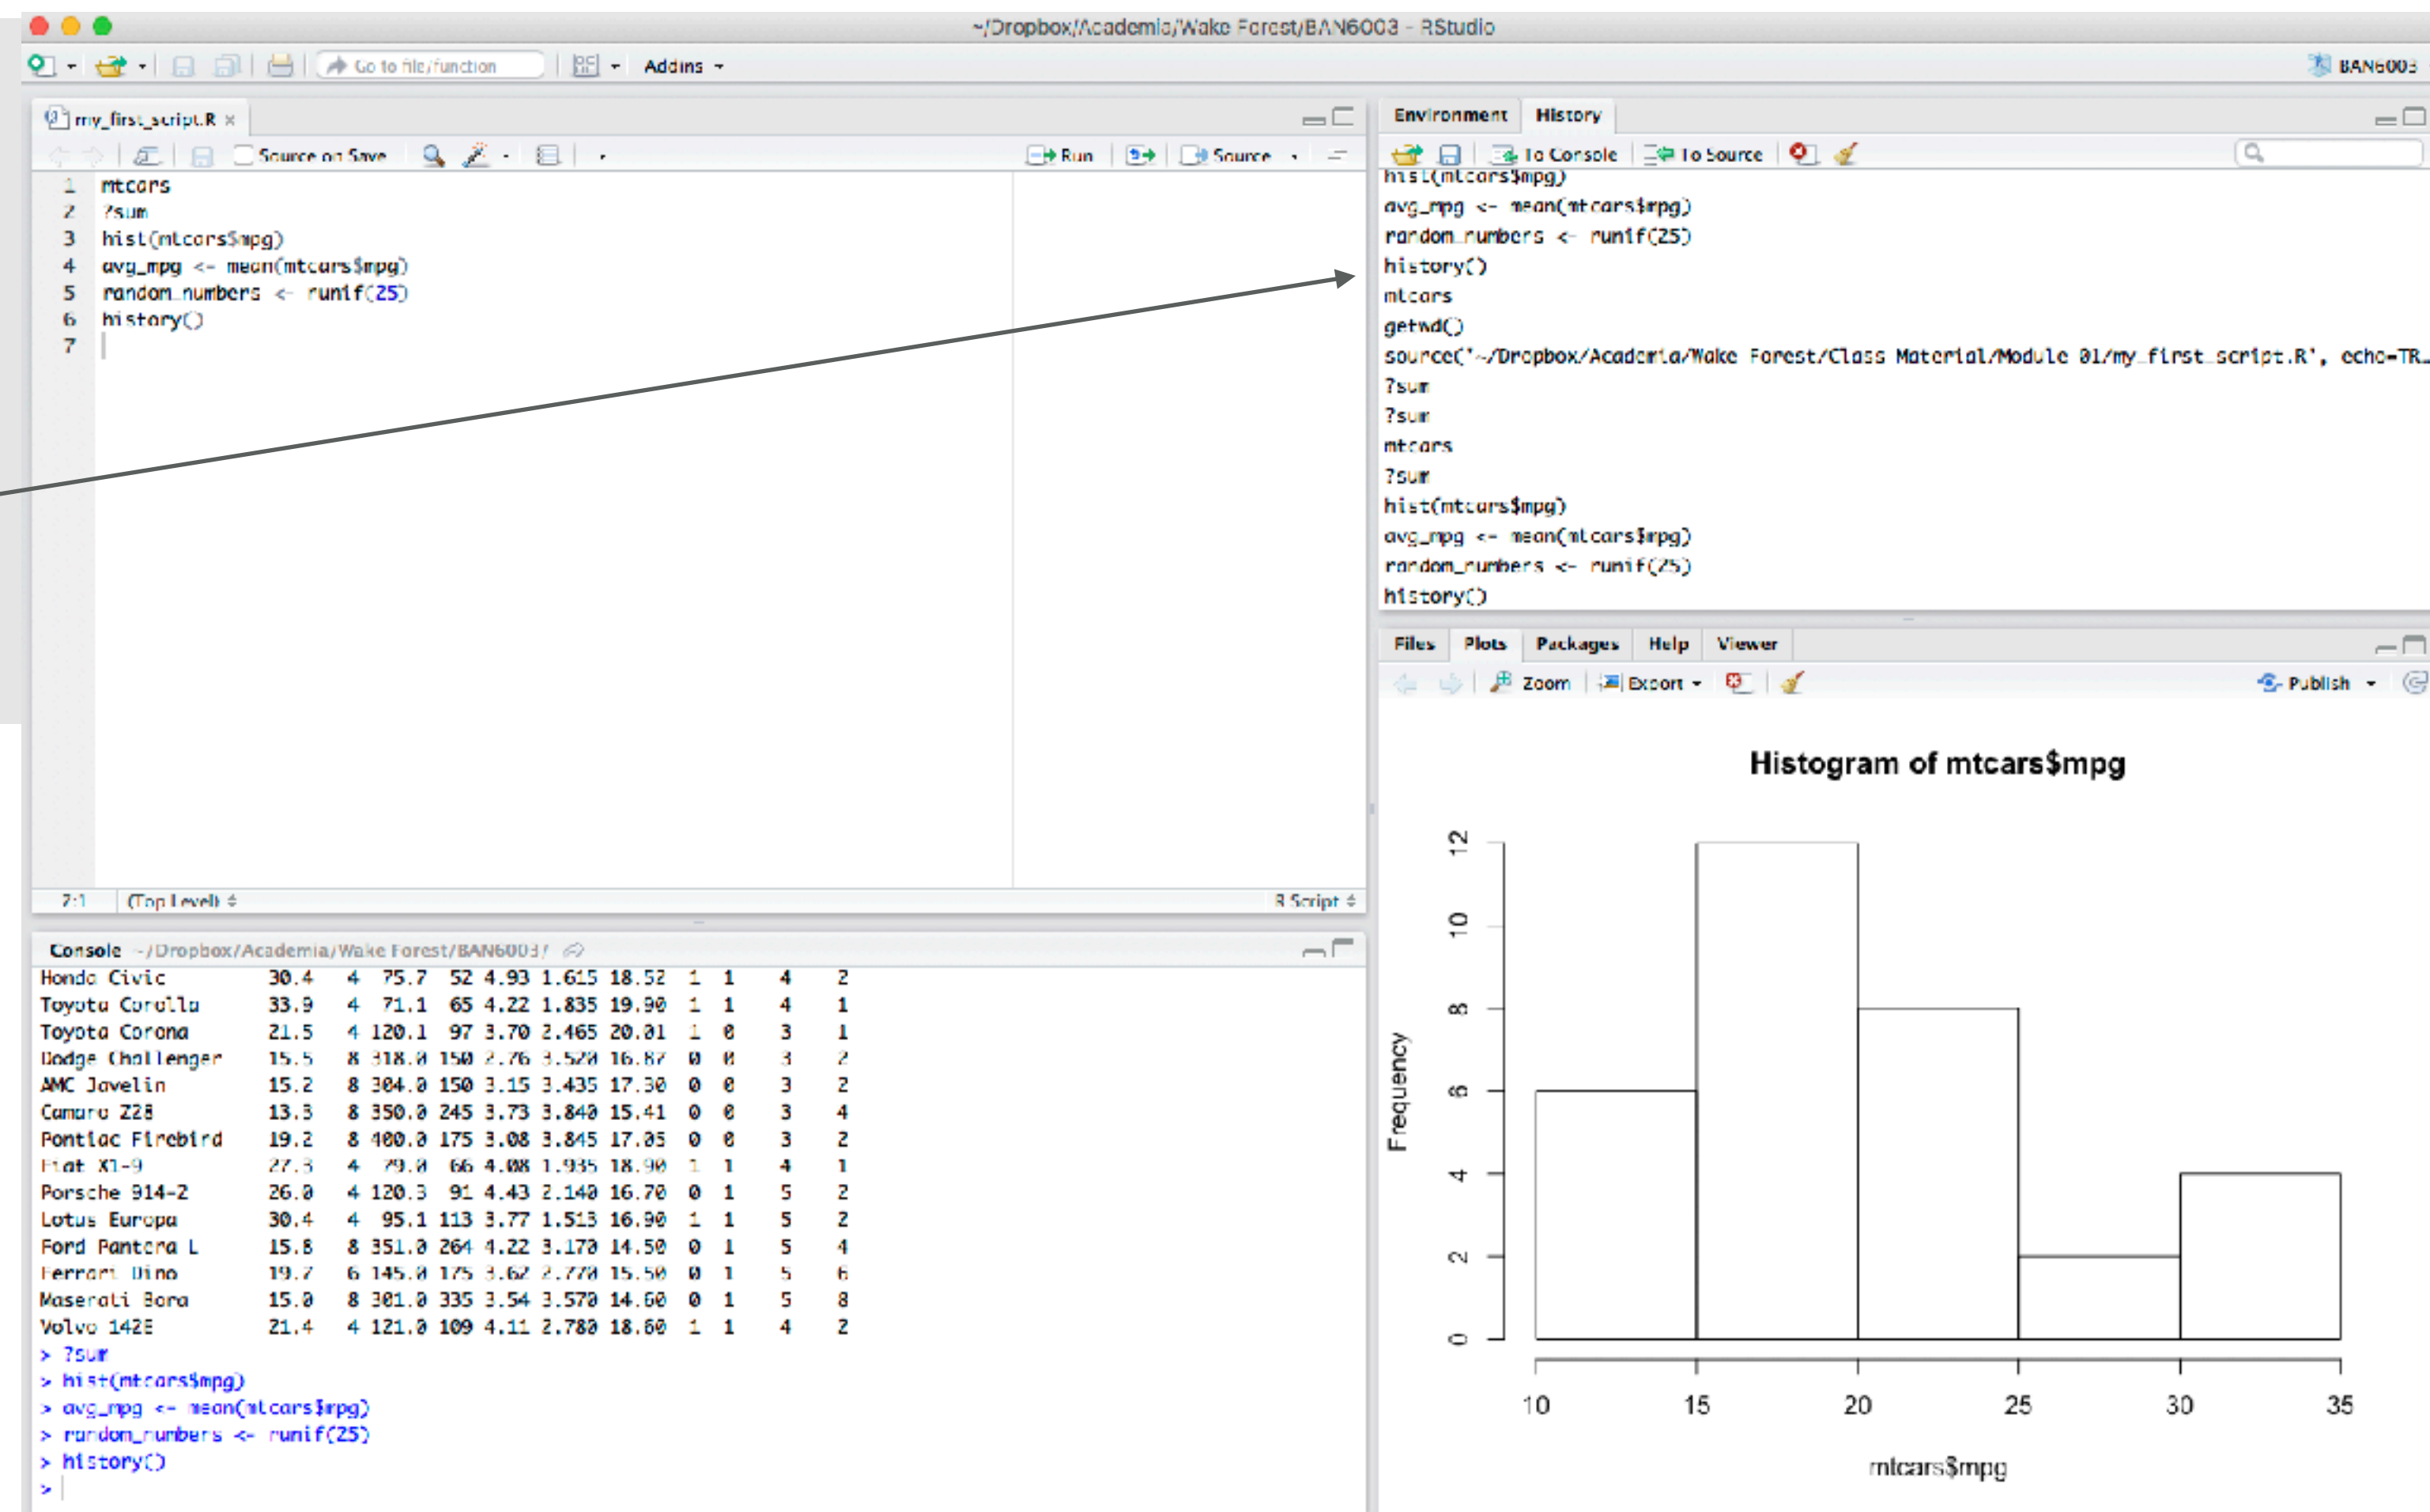
?sum

hist(mtcars\$mpg)

avg_mpg <- mean(mtcars\$mpg)

random_numbers <- runif(25)

history()



GETTING HELP

```
# provides details for specific function  
help(sqrt)
```

```
# provides same information as help(functionname)  
?sqrt
```

```
# provides examples for said function  
example(sqrt)
```

External to R:

Google: just add "with R" at the end of any search.

Stack Overflow: a searchable Q&A site oriented toward programming issues. 75% of my answers come from SO

Cross Validated: a searchable Q&A site oriented toward statistical analysis.

R-bloggers: a central hub of content from over 500 bloggers who provide news and tutorials about R.

SET YOUR WORKING DIRECTORY

```
# get your current working directory
```

```
getwd()
```

```
[1] "/Users/bradleyboehmke/Dropbox/Academia/Wake Forest/BAN6003"
```

```
# set your working directory
```

```
setwd("/Users/bradleyboehmke/Dropbox/Academia/Wake Forest")
```

```
getwd()
```

```
[1] "/Users/bradleyboehmke/Dropbox/Academia/Wake Forest"
```

Keeping your files organized is critical

YOUR TURN!

*Set your working directory to the “Intro to R Bootcamp”
folder you downloaded for this class.*

R AS A CALCULATOR

```
# Uses PEMDAS convention for order of operations
```

```
4 + 3 / 10 ^ 2
```

```
## [1] 4.03
```

```
4 + (3 / 10 ^ 2)
```

```
## [1] 4.03
```

```
(4 + 3) / 10 ^ 2
```

```
## [1] 0.07
```

```
# large/small numbers will be displayed in scientific notation
```

```
1 / 17 ^ 7
```

```
## [1] 2.437011e-09
```

```
# Undefined calculations result in Inf or NaN
```

```
1 / 0
```

```
## [1] Inf
```

```
Inf - Inf
```

```
## [1] NaN
```


THE ASSIGNMENT (<-) OPERATOR

```
x <- 3      # GOOD  
x = 3      # BAD
```

```
# we can increment (build onto) existing objects
```

```
x <- x + 1
```

```
x
```

```
## [1] 4
```

```
# must be specific
```

```
X
```

```
Error: object 'x' not found
```

YOUR TURN!

Economic Order Quantity Model:

$$Q = \sqrt{\frac{2DK}{h}}$$

*Calculate **Q** where:*

$$D = 1000$$

$$K = 5$$

$$h = 0.25$$

hint: `sqrt(x)` = \sqrt{x}

SOLUTION

```
D <- 1000
```

```
K <- 5
```

```
h <- .25
```

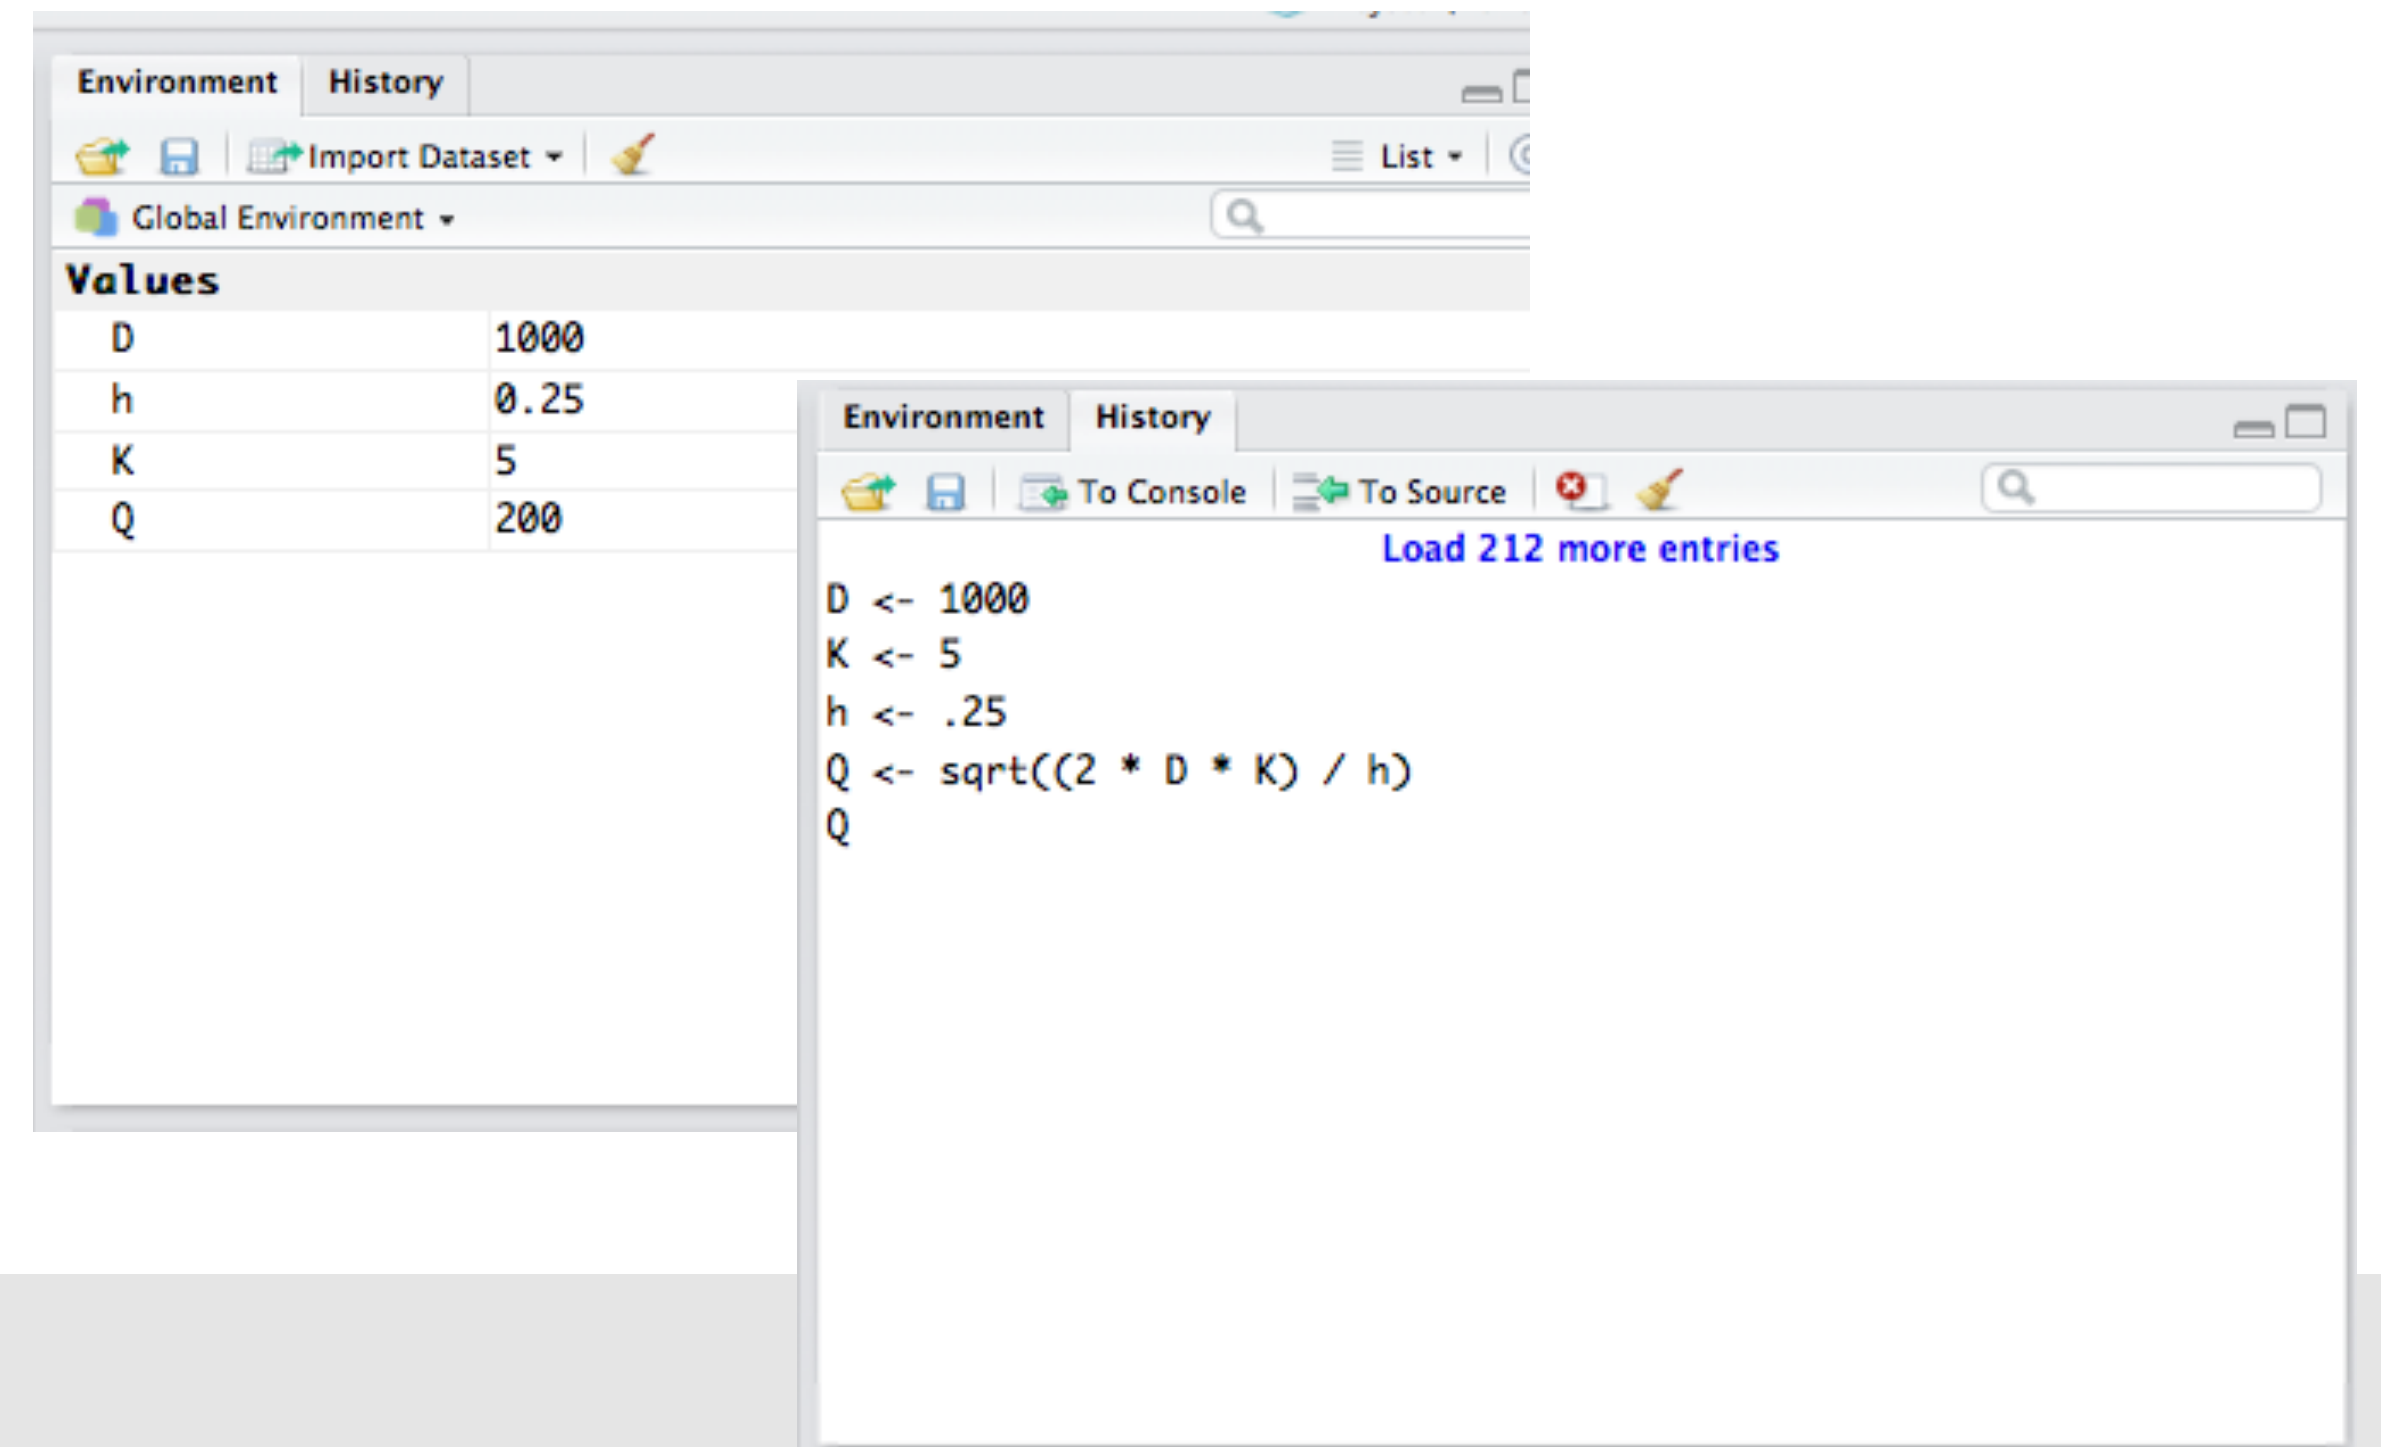
```
Q <- sqrt((2 * D * K) / h)
```

```
Q
```

```
## [1] 200
```


WORKSPACE ENVIRONMENT

- You should now have 4 objects in your global environment
- History tab will show your recent code



To list and remove objects in your global environment:

```
# list all objects
ls()
## [1] "D" "h" "K" "Q"
```

```
# remove defined object from the environment
rm(D)
```

```
# removes everything in the working environment -- use with caution!
rm(list = ls())
```

PACKAGES

The fundamental unit of shareable code is the package.

CRAN: 10,000+

Bioconductor: 1,000+

GitHub: Many more plus beta versions for updated packages not yet published

So how do we install these packages?

```
# install packages from CRAN
install.packages("packagename")
```

```
# install packages from Bioconductor
source("http://bioconductor.org/biocLite.R")
biocLite()
biocLite("packagename")
```

```
# install packages from GitHub
install.packages("devtools")
devtools::install_github("username/packagename")
```

```
# only required the first time
# only required the first time
```

```
# only required the first time
```

YOUR TURN!

Download these packages from CRAN:

tidyverse

nycflights13

SOLUTION

```
install.packages("tidyverse")  
install.packages("nycflights13")
```

```
# alternative  
install.packages(c("tidyverse", "nycflights13"))
```

For a full list of useful packages see this guide: <http://bit.ly/lx9vkzV>

PACKAGES

Loading packages:

```
# load the package to use in the current R session  
library(tidyverse)
```

```
# use a particular function within a package without loading the package  
stringr::str_replace()
```

Getting help on packages:

```
# provides details regarding contents of a package  
help(package = "tidyr")
```

```
# list vignettes available for a specific package  
vignette(package = "tidyr")
```

```
# view specific vignette  
vignette("tidy-data")
```

WHAT IS TIDYVERSE?



WHAT TO REMEMBER



FUNCTIONS TO REMEMBER

Operator/Function	Description
<code>help()</code> , <code>?</code> , <code>example()</code>	Get help on functions and provide examples
<code>getwd()</code> , <code>setwd()</code>	Get and set your working directory
<code>+</code> , <code>-</code> , <code>*</code> , <code>/</code> , <code>^</code>	Arithmetic
<code><-</code>	Assignment operator
<code>ls()</code> , <code>rm()</code>	list and remove objects in your global environment
<code>install.packages()</code> , <code>library()</code>	Install and load packages
<code>vignette()</code>	View/list package vignette

